
THE JOURNEY FROM WILD TO TEXTBOOK DATA: A CASE STUDY FROM THE NATIONAL LONGITUDINAL SURVEY OF YOUTH

A PREPRINT

Dewi Amaliah

Department of Econometrics and Business Statistics
Monash University
Clayton, VIC 3800
dama0007@student.monash.edu

Dianne Cook

Department of Econometrics and Business Statistics
Monash University
Clayton, VIC 3800
dicook@monash.edu

Emi Tanaka

Department of Econometrics and Business Statistics
Monash University
Clayton, VIC 3800
emi.tanaka@monash.edu

Kate Hyde

Department of Econometrics and Business Statistics
Monash University
Clayton, VIC 3800

Nicholas Tierney

September 26, 2021

Abstract

The National Longitudinal Survey of Youth (NLSY79) is a prominent open data source that has been important for educational purposes and multidisciplinary research on longitudinal data. Subsets of this data can be found in numerous textbooks and research articles. However, the steps and decisions taken to get from the raw data to the textbook data is never clearly articulated. This article describes our journey when trying to re-create a textbook example data set from the original database, with the goal being to refresh the textbook data more regularly. Thus, this paper demonstrates the process – extracting, tidying, cleaning, and exploring with documentation – to make refreshed data available for education or research. Three new data sets, and the code to produce them are provided in an accompanying open source R package, called **yowie**. As a result of this process, some recommendations are also made for the NLSY79 curators for incorporating data quality checks, and providing more convenient samples of the data to potential users.

Keywords Data cleaning; Data tidying; Longitudinal data; NLSY79; Open data; Initial data analysis; Outlier detection; Robust linear regression

1 Introduction

“Open data” is data that is freely accessible, modifiable, and shareable by anyone for any purpose (Open Knowledge Foundation 2021). This type of data can be useful as example data in statistical textbooks and for research purposes. However, open data is often referred to as what we might call “wild data” because it requires substantial cleaning and tidying to tame it into textbook shape. Huebner Marianne, Vach, and Cessie (2016) emphasize that making the data cleaning process accountable and transparent is imperative and essential for the integrity of downstream statistical analyses and model building (M. Huebner et al. 2020). Data cleaning can be considered to be a part of what is called “initial data analysis” (IDA) (Chatfield 1985). In IDA one would also explore the data, especially to check if the data is consistent with assumptions required for modeling. This is also related to exploratory data analysis (EDA), coined by Tukey (1977) with a focus on learning from data. EDA can be considered to encompass IDA. Dasu and Johnson (2003) say that data cleaning and exploration, without naming it as IDA, is a difficult task and consumes 80% of the data mining task.

Despite its importance, this IDA stage is often undervalued and neglected (Chatfield 1985). There are few research papers that document the data cleaning (Wickham 2014). Furthermore, the decisions made in this stage often go unreported in the sense that IDA is often performed in an unplanned and unstructured way and is only shared among restricted parties (M. Huebner et al. 2020).

This paper demonstrates the steps of IDA, and documents the process, for a prominent open data source, National Longitudinal Survey of Youth 1979, henceforth referred to NLSY79. This data has been playing an important role in research in various disciplines, including but are not limited to economics, sociology, education, public policy, and public health for more than a quarter of the century (Pergamit et al. 2001). In addition, this is considered a carefully designed longitudinal survey with high retention rates, making it suitable for life course research (Pergamit et al. 2001; Cooksey 2017). According to Cooksey (2017), thousands of articles, and hundreds of book chapters and monographs have utilized this data. Moreover, the NLSY79 is considered the most widely used and most important cohort in the survey data (Pergamit et al. 2001).

Singer and Willett (2003) used the wages and other variables of high school dropouts from the NLSY79 data as an example data set to illustrate longitudinal data modeling of wages on workforce experience, with covariates education and race. Our aim is to refresh this textbook data with data from 1994 through to the latest data reported in 2018. Here, we investigate the process of getting from the raw NLSY79 data to a textbook data set as similar as possible to that provided by Singer and Willett (2003). We should also note that race is a variable in the original data set, and for compatibility it is also provided with the refreshed data, for the purposes of studying racism, not race (Fullilove 1998),

Our process of cleaning builds heavily on the tidyverse approach (Wickham et al. 2019). The data is first organised into “tidy data” (Wickham 2014) and then further wrangled using the data pipeline and split-apply-combine approach (Wickham 2011). The resulting data is provided in a new R package called `yowie` which includes the code so that the process is reproducible, and could be used to further refresh the data as new records are made available in the NLSY79 database.

This paper is structured in the following way. Section 2 describes the NLSY79 data source. Section 3 presents the steps of cleaning the data, including getting and tidying the data from the NLSY79 and IDA to find and repair anomalies. Our final subset is compared the old textbook subset in Section 4. Finally, Section 5 summarizes the contribution and makes recommendations for the NLSY79 data curators.

2 The NLSY79

2.1 Database

The NLSY79 is a longitudinal survey held by the U.S Bureau of Labor Statistics that follows the lives of a sample of American youth born between 1957-1964 (The U.S. Bureau of Labor Statistics, n.d.). The cohort originally included 12,686 respondents aged 14-22 when first interviewed in 1979. It comprised of Blacks, Hispanics, economically disadvantaged non-Black non-Hispanics, and youth in the military. In 1984 and 1990, two sub-samples were dropped from the interview; the dropped subjects are the 1,079 members of the military sample and 1,643 members of the economically disadvantaged non-Black non-Hispanics, respectively. Hence, 9,964 respondents remain in the eligible samples. The surveys were conducted annually from 1979 to

1994 and biennially thereafter. Data are currently available from Round 1 (1979 survey year) to Round 28 (2018 survey year).

Although the main focus area of the NLSY is labor and employment, the NLSY also covers several other topics, including education; training and achievement; household, geography, and contextual variables; dating, marriage, cohabitation; sexual activity, pregnancy, and fertility; children; income, assets and program participation; health; attitudes and expectations; and crime and substance use.

There are two ways to conduct the interview of the NLSY79, which are face-to-face interviews or by telephone. In recent survey years, more than 90 percent of respondents were interviewed by telephone (Cooksey 2017).

2.2 Target data

This paper aims, as mentioned in Section 1, to refresh the NLSY79 data used in Singer and Willett (2003). This data set contains the information of yearly mean hourly wages from the NLSY79 subcohort, with education and race as covariates. The subcohort correspond to male high-school dropouts that participated in the study from 14 to 17 years old. Accordingly, to refresh this data, we make use of NLSY79 cohort’s demographic profiles and wages data of people who completed education up to high school.

Since no specific criteria on high school dropouts was defined in Singer and Willett (2003), we define high school dropouts as respondents who only completed either 9th to 11th grade, or who completed 12th when their age is at least 19 years old (meaning that the individual may have dropped out of high school but returned at a later time to complete high school education).


3 Data cleaning

Loo and Jonge (2018) describe the notion of a “statistical value chain” where the production stages of data cleaning process are earmarked as raw data (data from the data source), input data (data organised with correct type and identified variables) and valid data (data that faithfully represent the variables). In this section, we outline the steps to download the raw data (Section 3.1) and then tidying the raw data into input data, specifically for wages data (Section 3.1.1), demographic variables (Section 3.1.2) and the employment variables (Section 3.1.3), so that the resulting input data can be used downstream for validating the data as described in Section 3.2.

3.1 Getting the data

The NLSY79 data contains a large scope of variables but for our aim, the scope required is limited to demographic profiles and wages data. More specifically, we went to the NLSY79 database website at <https://www.nlsinfo.org/content/cohorts/nlsy79/get-data>, clicked on the direct link to NLSY79 data and navigated as described below.

Navigating the data source to download the raw data

 NLSY79 (<https://www.nlsinfo.org/investigator/pages/search?s=NLSY79>)

✓ The CASEID will be always be selected.

✓ The 3 recommended demographic variable (sample ID, race and sex) were selected.

For the remaining variables, we went to the “Variable Search” tab and select variables as follows

▷ Education, Training and Achievement Scores

▷ Education ▷ Summary measures ▷ All schools ▷ By year ▷ Highest grade completed

✓ All 80 variables in Highest grade completed were selected.

▷ Employment

▷ Summary measures ▷ By job

▷ Hours worked

✓ All 442 primary variables in Hours worked were selected.


▷ Hourly wages

✓ All 156 variables in Hourly wages were selected.

▷ Household, Geography and Contextual Variables

▷ Context ▷ Summary measures ▷ Basic demographics ▷ Date of birth

✓ All 4 variables in Date of birth were selected.

 To download all 686 variables selected, we then navigate to the tab "Save / Download" then select the tab "Advanced Download". We select the R Source code, Codebook of selected variables, Short Description File and Comma-delimited datafile of selected variables with Reference Number as column headers. We name the filename "NLSY79" and press the download button.

The downloaded data set came as a zip file when extracted contains the following set of files:

- NLSY79.csv: csv format of the response data.
- NLSY79.dat: dat format of the response data.
- NLSY79.NLSY79: Tagset of variables that can be uploaded to the website to recreate the data set.
- NLSY79.R: R script for reading the data into R and converting the variables' names and label into something more sensible. We utilize this code to do an initial data tidying which produces two data set, `categories_qnames` (where the observations are stored in categorical/interval values) and `new_data_qnames` (the observations are stored in integer form).

3.1.1 Tidying the wages data

The raw data, `new_data_qnames`, is organised such that each row corresponds to an individual. As respondents can have multiple jobs at specific years, the column names, such as `HRP1_1979`, `HRP2_1979`, `HRP1_1980` and `HRP2_1980`, contain the information about the job number up to 5 (`HRP1` = job 1, `HRP2` = job 2) and the year. The raw data consequently has a large number of columns (686 to be specific). The values in the cell under the variables that begin with `HRP` correspond to the hourly wage in dollars. A glimpse of this data output is shown below.

```
#> 'data.frame': 12686 obs. of 686 variables:
#> $ CASEID_1979 : int 1 2 3 4 5 6 7 8 9 10 ...
#> $ HRP1_1979 : int 328 385 365 NA 310 NA NA NA 214 NA ...
#> $ HRP2_1979 : int NA NA NA NA 375 NA NA NA NA NA ...
#> $ HRP3_1979 : int NA NA 275 NA NA NA NA NA NA NA ...
#> $ HRP4_1979 : int NA NA NA NA NA 250 NA NA NA NA ...
#> $ HRP5_1979 : int NA NA NA NA NA NA NA NA NA NA ...
#> $ HRP1_1980 : int NA 457 397 NA 333 275 300 394 200 318 ...
#> $ HRP2_1980 : int NA NA 367 NA NA NA NA NA NA NA ...
#> $ HRP3_1980 : int NA NA 380 NA NA NA 290 NA NA NA ...
#> $ HRP4_1980 : int NA NA NA NA NA NA NA NA NA NA ...
```

```
#> [list output truncated]
```

According to Wickham (2014), tidy data sets comply with three rules: (i) each variable forms a column, (ii) each observation forms a row, and (iii) each type of observational unit forms a table. The raw data does not comply the rules of the tidy data, thus we wrangle the data in form required by making use of the **tidyverse** suite of packages (Wickham et al. 2019), which include the main set of packages needed, namely **tidyr** (Wickham 2020b), **dplyr** (Wickham et al. 2020), and **stringr** (Wickham 2019). The tidy data form, where we have a variable that indicates the individual ID, year, job number, wage in dollars and demographic variables, is important for downstream analysis, such as visualisation and modelling that will be performed later.

[REMOVE] Unfortunately, the **new_data_qnames** did not meet these requirements in the way that data for each year and job for the same respondent is stored as one variable. Hence, the data contains a huge number of columns (686 columns). The example of the untidy data set is displayed in Table ???. The table intended to display the hourly rate of each respondent by job (HRP1 to HRP5) and by year (1979 and 1980). The table implies that the column headers are values of the year and job, not variable names. Consequently, the data should be tidied and wrangled first to extract the demographic and employment variables we want to put in the final data set. We mainly used **tidyr** (Wickham 2020b), **dplyr** (Wickham et al. 2020), and **stringr** (Wickham 2019) to do this job.

3.1.2 Tidying demographic variables

Age in 1979, **gender**, **race**, **highest grade completed**, and the **year when the highest grade completed** are the variables that we want to use in the cleaned data set.

Age in 1979 are derived from year and month of birth variables in the raw data. The variables have two versions, which are the 1979 version and the 1981 version. We only used the 1979 data and did a consistency check of those years and flag the inconsistent observations.

For **gender** and **race**, we only rename the column names. It is worth noting that using these two variable needs special attention. Gender, as reported in the data, only has two categories, which is recognised today as inadequate. Gender is not binary. Further, race is as reported in the database. When doing analysis with this variable, one should keep in mind that the purpose is to study racism rather than race.

The next step is tidying the **highest grade completed** variable. This variable came with several version in each year. We choose the revised May data because it seemed to have less missing and presumably has been checked. However, there is no revised May data for 2012, 2014, 2016, and 2018. Thus, we just use the ordinary May data for these years. The code for tidying this variable is provided in the supplementary materials.

Further, **highest grade completed** is measured and could be updated in each period of the survey. Hence, we only used the highest grade completed ever and derived the year when it is completed. The code for tidying this variable is also available in the supplementary materials.

Finally, we get all of the demographic profiles of the NLSY79 cohort. We then save this data as **demog_nlsy79**.

3.1.3 Tidying employment variables

The employment data comprises of three variables, i.e., **total hours of work per week**, **number of jobs that an individual has**, and **mean hourly wage**.

hours worked per week initially has only one version per job, no choice from 1979 to 1987 (QES-52A). From 1988 onward, when we had more options, we chose the variable for total hours, including time spent working from home (QES-52D). However, in 1993, this variable did not have all the five D variables (the first one and the last one were missing), so we used QES-52A variable instead. In addition, 2008 only had jobs 1-4 for the QES-52D variable (whereas the other years had 1-5), so we use these.

The same way was also deployed to tidy the rate of wage by year and by ID. The difference is that the hourly rate has only one version of each year. The hours of work and the hourly rate are then joined to calculate the number of jobs that a respondent has and their mean hourly wage. Some observations have 0 in their hourly rate and work hours, which is considered an invalid value. Thus, these observations are set to be NA. Besides, there are some observations with unusual hours of work. Here, we also censored the observations with > 84 hours of work a week to be NA.

Since our ultimate goal is to calculate the mean hourly wage, the number of jobs is calculated based on the availability of the `rate_per_hour` information. For example, the number of jobs of ID 1, based on `hours_work`, is 2. However, since the information of hourly rate of `job_02` is not available, the number of jobs is considered as 1.

Further, we calculated the mean hourly wage for each ID in each year using a weighted mean with the hours of work as the weight. However, there is a lot of missing value in `hours_work` variable. In that case, we only calculated the mean hourly wage based on the arithmetic/regular mean method. Hence, we created a new variable to flag whether the mean hourly wage is weighted or regular. Additionally, if an ID only had one job, we directly used their hourly wages information and flagged it as arithmetic mean.

```
#> # A tibble: 10 x 6
#>       id year mean_hourly_wage total_hours number_of_jobs is_wm
#>   <int> <dbl>         <dbl>         <int>         <dbl> <lgl>
#> 1     1   1979           3.28             38             1 FALSE
#> 2     1   1981           3.61             NA             1 FALSE
#> 3     2   1979           3.85             35             1 FALSE
#> 4     2   1980           4.57             NA             1 FALSE
#> 5     2   1981           5.14             NA             1 FALSE
#> 6     2   1982           5.71             35             1 FALSE
#> 7     2   1983           5.71             NA             1 FALSE
#> 8     2   1984           5.14             NA             1 FALSE
#> 9     2   1985           7.71             NA             1 FALSE
#> 10    2   1986           7.69             NA             1 FALSE
```

The employment and demographic variables are then joined. We also filtered the data to only have the cohort who completed the education up to 12th grade and participated at least five rounds in the survey and save it to an object called `wages`.

3.2 Initial data analysis

According to Huebner Marianne, Vach, and Cessie (2016), initial data analysis (IDA) is the step of inspecting and screening the data after being collected to ensure that the data is clean, valid, and ready to be deployed in the later formal statistical analysis. Moreover, Chatfield (1985) argues that the two main objectives of IDA are data description, which is to assess the structure and the quality of the data, and model formulation without any formal statistical inference.

In this paper, we conduct an IDA or a preliminary data analysis to assess the consistency of the data with the cohort information that the NLSY provides. In addition, we also aim to find the anomaly in the wages values using this approach. We mainly use graphical summaries to do the IDA using `ggplot2` (Wickham 2016) and `brlrgar` (Tierney, Cook, and Prvan 2020).

As stated previously, the respondents' ages ranged from 12 to 22 when first interviewed in 1979. Hence, we validate whether all of the respondents were in this range. Additionally, the NLSY also provides the number of the survey cohort by their gender (6,403 males and 6,283 females) and race (7,510 Non-Black/Non-Hispanic; 3,174 Black; 2,002 Hispanic). To validate this, we used the `demog_nlsy79`, i.e., the data with the survey years 1979 sample. Table 1 and Table 2 suggest that the demographic data we had is consistent with the sample information in the database.

The next step is that explore the mean hourly wage data. In this case, we only explore the wages data of up to 12th grade. Here, we use visualisation techniques to do IDA.

We take 36 samples randomly from the data and plot them, as shown in Figure 1. It implies that these respondents have a lot of variability in wages, for example, the IDs in panel numbers 5, 7, and 11. The plot also implies that the samples have a different pattern of mean hourly wages. Some have had flat wages for years but had a sudden increase in one particular year, then it gone down again, while the others experienced an upsurge in their wage, for instance, the IDs in panel 9.

The anomalies are also found in the total hours of work, where some observations work for 420 hours a week in total. According to Pergamit et al. (2001), one of the flaws of the NLSY79 employment data is that the NLSY79 collects the information of the working hours since the last interview. Thus, it might be challenging for the respondents to track the within-job hours' changes between survey years, especially for the respondents with fluctuated working hours or seasonal jobs. It even has been more challenging since

Table 1: The frequency table of the age at the start of the survey (?CHECK) in the full NSLY79 data

Age	Number of individuals
15	1,265
16	1,550
17	1,600
18	1,530
19	1,662
20	1,722
21	1,677
22	1,680

Table 2: The contingency table for gender and race for the full NLSY79 data.

Gender	Race			Total
	Hispanic	Black	Non-Black, Non-Hispanic	
Male	1,000 (15.62%)	1,613 (25.19%)	3,790 (59.19%)	6,403 (100.00%)
Female	1,002 (15.95%)	1,561 (24.84%)	3,720 (59.21%)	6,283 (100.00%)
Total	2,002 (15.78%)	3,174 (25.02%)	7,510 (59.20%)	12,686 (100.00%)

1994, where the respondents had to recall two years periods. This shortcoming might also contribute to the fluctuation of one’s wages data.

3.2.1 Replacing extreme values

As part of the IDA, which is the model formulation, we build a robust linear regression model to treat the extreme values in the data. Robust linear regression yields an estimation robust to the influence of noise or contamination (Koller 2016). It also aims to detect the contamination by weighting each observation based on how “well-behaved” they are, known as robustness weight. Observations with lower robustness weight are suggested as an outlier by this method (Koller 2016).

Since we work with longitudinal data, we build the model for each ID instead of the overall data. The robust mixed model could be the best model to be employed in this case. However, this method is too computationally and memory expensive, especially for a large data set, like the NLSY79 data. Thus, the model for each ID is built utilizing the `nest` and `map` function from `tidyr` (Wickham 2020b) and `purrr` (Henry and Wickham 2020), respectively.

We build the model using the `rlm` function from `MASS` package (Venables and Ripley 2002). We set the `mean_hourly_wage` and `year` as the dependent and predictor, respectively. Furthermore, we use M-Estimation with Huber weighting, where the observation with a slight residual gets a weight of 1, while the larger the residual, the smaller the weight (less than 1) (UCLA: Statistical Consulting Group 2021). However, the challenging part of detecting the anomaly using the robustness weight is determining the weight threshold in which the observations are considered outliers. Moreover, it should be noted that not all the outliers are due to an error. Instead, it might be that one had reasonably increasing or decreasing wages in a particular period.

To minimize the risk of mistakenly identifying an outlier as an “erroneous outlier,” we simulate some thresholds and study how they affect the data. We find that 0.12 is the most reasonable value to be the threshold to minimize that drawback’s risk because it still captures the sensible spikes in the data. In other words, we keep maintaining the natural variability of the wages while minimizing anomalies because of the error in the data recording. After deciding the threshold, we impute the observations whose weights are less than 0.12 with the models’ predicted value. We then flag those observations in a new variable called `is_pred`.

Figure 3 shows the mean hourly wage before and after the extreme values are replaced. It implies that the fluctuation can still be observed in the data after the treatment. However, the large spikes, which are considered “erroneous outliers,” are already eliminated from the data. Hence, the model produces a data set with a more reasonable degree of fluctuation.

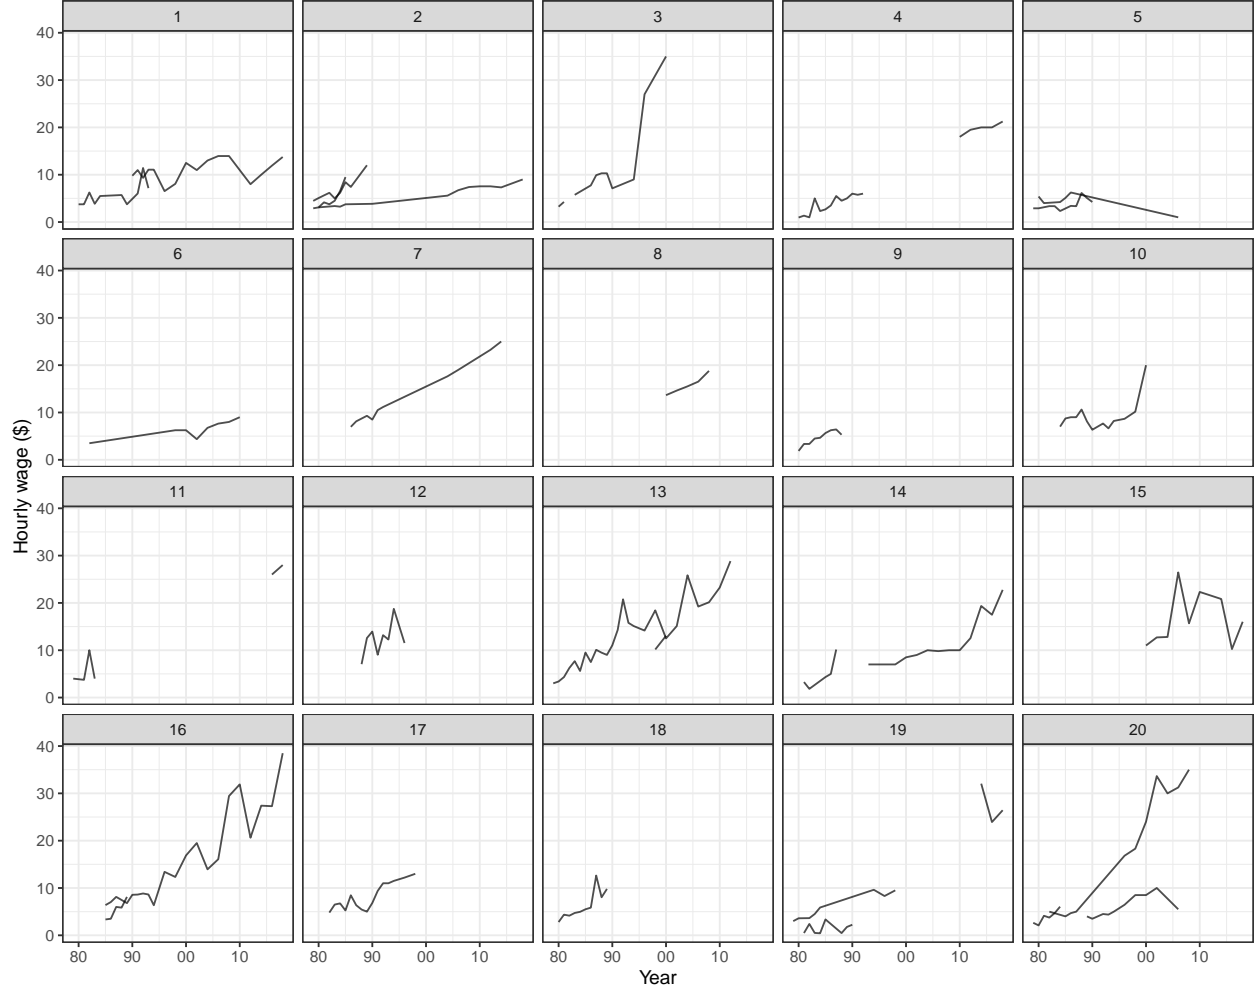


Figure 1: Longitudinal profiles of wages for a random sample of 20 individuals in the refreshed data. Most, but not all, individuals here experienced increasing wages over time, and several have experienced considerable fluctuation in wages. Some individuals are only measured for a short period.

Further, 4 A shows that after eliminating the extreme values, the highest value has decreased to be around \$350. The spikes are still observed but not as extreme as the original data set. In Figure 4 B, we plot the three features of mean hourly wages, namely the minimum, median, and maximum value. We still can see some extreme values in maximum wages, but consider it as a natural variability of the data. In 4 C, we can see that after imputing the extreme value in ID=39, we can see how the wages change over the years more clearly.

Finally, we save the imputed data and set the appropriate data type for the variables.

3.3 Recap

In this section, we have demonstrate a workflow from getting the data from the NSLY79 database until a ready-analysed data. Referring to Loo and Jonge (2018), this section shows part of statistical value chain, which is data cleaning from raw, input, and valid data. The raw demographic and employment data from the database are not in a tidy format. Hence, we perform data wrangling to get the input data. We inspect the input data to find eligible observations to the target data. One of the target data is wages data of respondents who completed up to 12th grade and participated at least in 5 rounds of survey. This needs us to join the extracted demographic and employment data.

Further, using IDA, we inspect on the existence of anomalies/extreme values in the data using the weight from robust linear model. We then replace these extreme values using their predicted value. From this stage




Figure 2: Summary plots to check the cleaned data reveal more cleaning is necessary: longitudinal profiles of wages for all individuals 1979-2018 (A), boxplots of minimum, median, and maximum wages of each individual (B), and one individual (id=39) with an unusual wage relative to their years of data (C). Some values of hourly wages are unbelievable, and some individuals have extremely unusual wages in some years.

forward, we assume that we already have valid data. We also subset the data to get the wages of high school dropouts. We then make these three data sets and the cleaning workflow documentation publicly available through an R data container package called *yowie*. A visualisation of these process is displayed in Figure 5.

4 Comparison of refreshed with the original data

Now is the time to see how close we have come to the original textbook data. The original set contains wages of high school dropouts wages (Singer and Willett 2003) from 1979 through to 1994. The refreshed data set for comparison is also wages on high school drop outs, over 1979 to 2018. The original set is available in the R package *brolgar* and the refreshed data is available in the R package *yowie*.

There are two aspects of the original data that make a direct comparison difficult. The time variable provided is “experience in the workforce.” It is not clear how this is calculated. One would expect that there is a record of the day the individual first started a job, and this is used to adjust the year of collection. We have not been able to find this information in the database. Secondly, wages were inflation-adjusted to 1990 prices, which is not done for the refreshed data.

Figure 6 shows the longitudinal profiles of the two sets. (A direct ID matching is not possible.) There are 888 individuals in the original and 1,188 individuals in the refreshed data. This suggests some individuals were removed in the original data. Overall, the profiles show some similarity.

5 Summary

This paper has performed stages to make open data suitable for textbook data or make it ready for research. In the first stage, we showed the steps performed to get the data from the NLSY79 database. Since the data format is untidy, we showed how the data had been tidied. After that, we conducted an initial data analysis to investigate and screen the quality of the data. We found and fixed the anomalous observations in the data set using the robust linear regression model with its predicted values. We also performed an example of exploratory data analysis using the cleaned data set.

This paper has also demonstrated data cleaning documentation by providing all of the codes in performing data tidying and initial data analysis. Thus, this paper provided an opportunity to continuously refresh the textbook data whenever the updated data is published in the NLSY79 database. It could be done by following the documentation of the code that is provided in this paper.

Moreover, the documentation also includes how we generated the robustness weight and how we decided the threshold of anomalous observations. It is also documented along with the flag of whether an observation

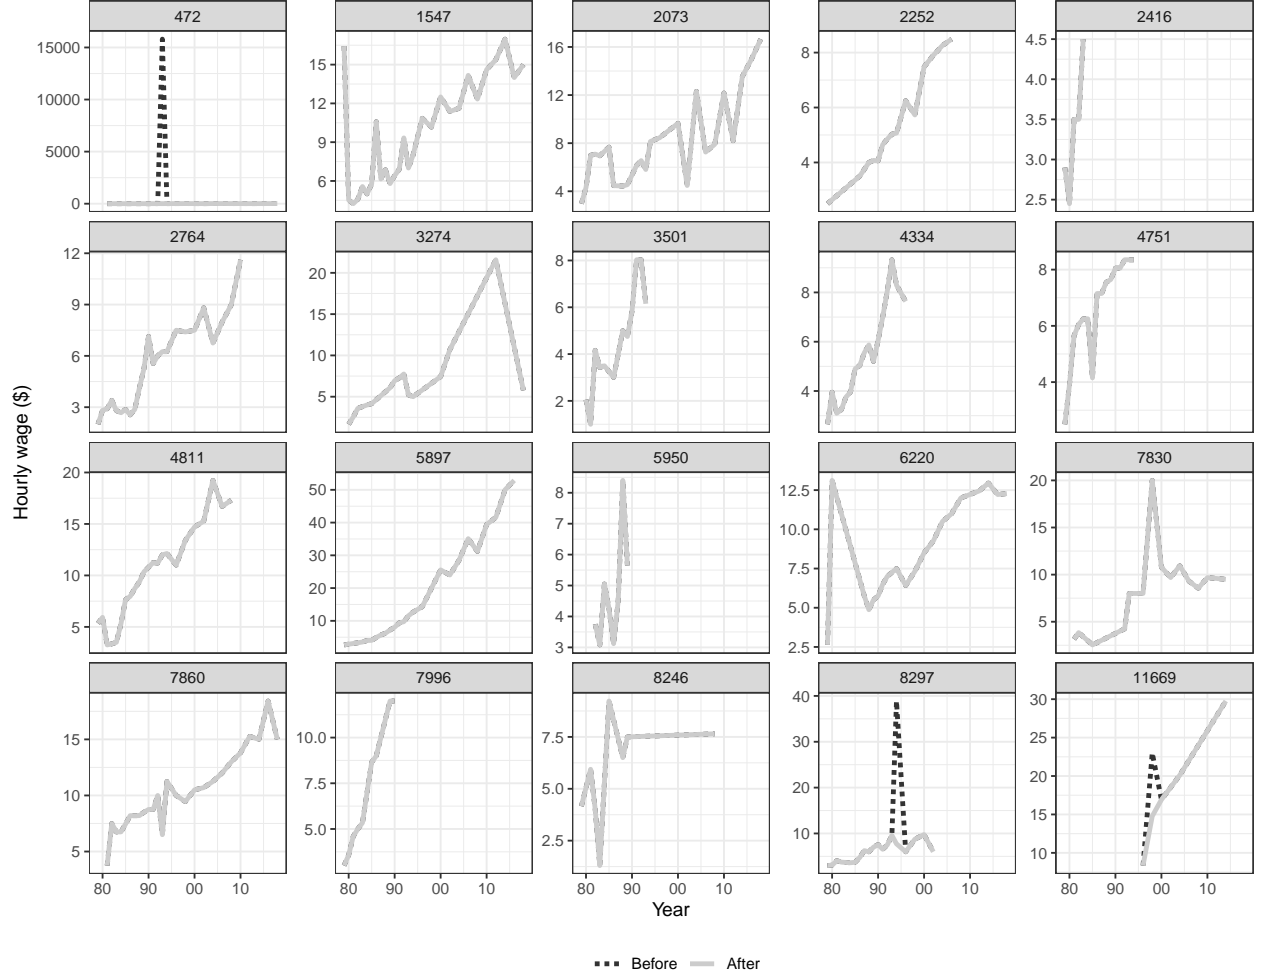


Figure 3: Comparison between the original (black dots) and the corrected (solid grey) mean hourly wage for a sample of individuals. A robust linear model prediction was used to correct mean hourly wages value. We can see that some extreme spikes, corresponding to implausible wages, have been replaced with values more similar to wages in neighboring years, but otherwise the profiles are not changed. Some spikes might remain when wage values are plausible.

is an imputed value or not. Accordingly, if somebody wishes to make another decision, it can be done by making a small change in the code provided. Further, this paper is also supplemented by a **shiny** (Chang et al. 2020) app as a simulation tool to customize the weight threshold.

Finally, this paper implies that data providers should design a database that is able to produce tidy data sets. A data provider should also check for data anomalies before the data publishing or at least provides a set of rules or threshold values. For example, in this case, is the threshold of reasonable wages. This will greatly support the data users to validate and set the same understanding of which data are considered outliers. Moreover, providing validation rules would facilitate any established data validation tool, such as **validate** (van der Loo and de Jonge 2021) package. In this case, we cannot use this handy validation package due to the absence of validation rules.

6 Acknowledgements

We would like to thank Aarathy Babu for the insight and discussion during the writing of this paper.

The entire analysis is conducted using R (R Core Team 2020) in **rstudio** using these packages: **tidyverse** (Wickham et al. 2019), **ggplot2** (Wickham 2016), **dplyr** (Wickham et al. 2020), **readr** (Wickham and

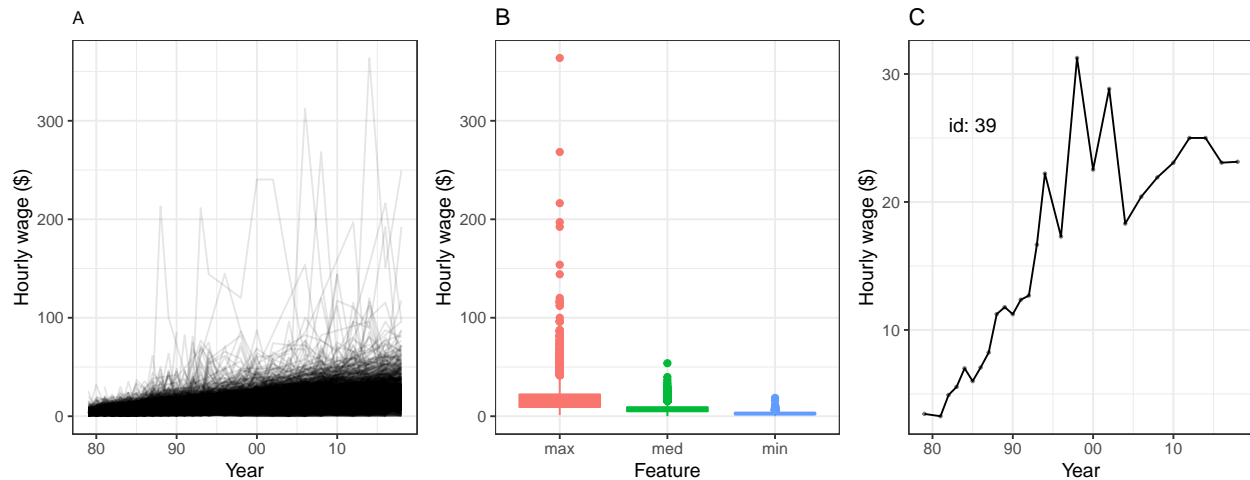


Figure 4: Re-make of the summary plots of the fully processed data suggest that it is now in a reasonable state: longitudinal profiles of wages for all individuals 1979-2018 (A), boxplots of minimum, median, and maximum wages of each individual (B), and one individual with an unusual wage relative to their years of data (C).

Hester 2020), `tidyr` (Wickham 2020b), `stringr` (Wickham 2019), `purrr` (Henry and Wickham 2020), `broom` (Robinson, Hayes, and Couch 2021), `brlgar` (Tierney, Cook, and Prvan 2020), `patchwork` (Pedersen 2020), `kableExtra` (Zhu 2019), `MASS` (Venables and Ripley 2002), `janitor` (Firke 2020), `DiagrammeR` (Iannone 2020), `rsvg` (Ooms 2020), `webshot` (Chang 2019), `mgcv` (Wood 2003), `tsibble` (Wang, Cook, and Hyndman 2020), and `modelr` (Wickham 2020a). The paper are generated using `knitr` (Xie 2014), `rmarkdown` (Xie, Dervieux, and Riederer 2020), and `rticles` (Allaire et al. 2021).

7 Supplementary Materials

- **Codes** : R script to reproduce data tidying and cleaning are available in this Github Repository.
- **R Package yowie**: yowie is a data container R package that contains 3 datasets, namely the high school mean hourly wage data, high school dropouts mean hourly wage data, and demographic data of the NLSY79 cohort. This package could be accessed here.
- **shiny app**: An web interactive shiny app to run a simulation to customize the weight threshold. This app could be accessed here.

References

- Allaire, JJ, Yihui Xie, R Foundation, Hadley Wickham, Journal of Statistical Software, Ramnath Vaidyanathan, Association for Computing Machinery, et al. 2021. *Rticles: Article Formats for r Markdown*. <https://CRAN.R-project.org/package=rticles>.
- Chang, Winston. 2019. *Webshot: Take Screenshots of Web Pages*. <https://CRAN.R-project.org/package=webshot>.
- Chang, Winston, Joe Cheng, JJ Allaire, Yihui Xie, and Jonathan McPherson. 2020. *Shiny: Web Application Framework for r*. <https://CRAN.R-project.org/package=shiny>.
- Chatfield, C. 1985. "The Initial Examination of Data." *Journal of the Royal Statistical Society. Series A. General* 148 (3): 214–53.
- Cooksey, Elizabeth C. 2017. "Using the National Longitudinal Surveys of Youth (NLSY) to Conduct Life Course Analyses." In *Handbook of Life Course Health Development*, edited by Richard M. Lerner Neal Halfon Christopher B. Forrest, 561–77. Cham: Springer. https://doi.org/https://doi.org/10.1007/978-3-319-47143-3_23.




Figure 5: The stages of data cleaning from the raw data to get three datasets contained yowie. n individuals means the number of respondents included in each stage, while n observations means the number of rows in the data. The color represented the stage of data cleaning in statistical value chain (van der Loo & de Jonge, 2018). Pink, blue, and green represents the raw, input, and valid data, respectively.




Figure 6: Comparison of original textbook example (A) with refreshed data (B). The original data was inflation-adjusted to 1990 prices and the individual’s time of collection was converted to a length of experience in the workforce, which make it difficult to precisely compare the two sets.

- Dasu, Tamraparni, and Theodore Johnson. 2003. *Exploratory Data Mining and Data Cleaning*. Wiley Series in Probability and Statistics. Hoboken: WILEY.
- Firke, Sam. 2020. *Janitor: Simple Tools for Examining and Cleaning Dirty Data*. <https://CRAN.R-project.org/package=janitor>.
- Fullilove, M. T. 1998. “Comment: Abandoning “Race” as a Variable in Public Health Research—an Idea Whose Time Has Come.” *American Journal of Public Health* 88 (9): 1297–98.
- Henry, Lionel, and Hadley Wickham. 2020. *Purrr: Functional Programming Tools*. <https://CRAN.R-project.org/package=purrr>.
- Huebner, Marianne, Werner Vach, and Saskia le Cessie. 2016. “A Systematic Approach to Initial Data Analysis Is Good Research Practice.” *The Journal of Thoracic and Cardiovascular Surgery* 151 (1): 25–27.
- Huebner, Marianne, Werner Vach, Saskia le Cessie, Carsten Oliver Schmidt, and Lara Lusa. 2020. “Hidden Analyses: A Review of Reporting Practice and Recommendations for More Transparent Reporting of Initial Data Analyses.” *BMC Medical Research Methodology* 20 (1): 61–61.
- Iannone, Richard. 2020. *DiagrammeR: Graph/Network Visualization*. <https://CRAN.R-project.org/package=DiagrammeR>.
- Koller, Manuel. 2016. “Robustlmm: An r Package for Robust Estimation of Linear Mixed-Effects Models.” *Journal of Statistical Software* 75 (6): 1–24.
- Loo, Mark van der, and Edwin de Jonge. 2018. *Statistical Data Cleaning with Applications in r*.
- Ooms, Jeroen. 2020. *Rsvg: Render SVG Images into PDF, PNG, PostScript, or Bitmap Arrays*. <https://CRAN.R-project.org/package=rsvg>.
- Open Knowledge Foundation. 2021. “Open Definition. Defining Open in Open Data, Open Content, and Open Knowledge.” 2021. <http://opendefinition.org/od/2.1/en/>.
- Pedersen, Thomas Lin. 2020. *Patchwork: The Composer of Plots*. <https://CRAN.R-project.org/package=patchwork>.

- Pergamit, Michael R., Charles R. Pierret, Donna S. Rothstein, and Jonathan R. Veum. 2001. "Data Watch: The National Longitudinal Surveys." *The Journal of Economic Perspectives* 15 (2): 239–53.
- R Core Team. 2020. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. <https://www.R-project.org/>.
- Robinson, David, Alex Hayes, and Simon Couch. 2021. *Broom: Convert Statistical Objects into Tidy Tibbles*. <https://CRAN.R-project.org/package=broom>.
- Singer, Judith D, and John B Willett. 2003. *Applied Longitudinal Data Analysis: Modeling Change and Event Occurrence*. Oxford u.a: Oxford Univ. Pr.
- The U.S. Bureau of Labor Statistics. n.d. "National Longitudinal Survey of Youth 1979." Available at <https://www.nlsinfo.org/content/cohorts/nlsy79> (2021/25/02).
- Tierney, Nicholas, Di Cook, and Tania Prvan. 2020. *Brolgar: BRowse over Longitudinal Data Graphically and Analytically in r*. <https://github.com/njtierney/brolgar>.
- Tukey, John W. (John Wilder). 1977. *Exploratory Data Analysis*. Addison-Wesley Series in Behavioral Science. Reading, Mass.: Addison-Wesley Pub. Co.
- UCLA: Statistical Consulting Group. 2021. "Robust Regression | r Data Analysis Examples." February 2021. <https://stats.idre.ucla.edu/r/dae/robust-regression/>.
- van der Loo, Mark P. J., and Edwin de Jonge. 2021. "Data Validation Infrastructure for R." *Journal of Statistical Software* 97 (10): 1–31. <https://doi.org/10.18637/jss.v097.i10>.
- Venables, W. N., and B. D. Ripley. 2002. *Modern Applied Statistics with s*. Fourth. New York: Springer. <http://www.stats.ox.ac.uk/pub/MASS4>.
- Wang, Earo, Dianne Cook, and Rob J Hyndman. 2020. "A New Tidy Data Structure to Support Exploration and Modeling of Temporal Data." *Journal of Computational and Graphical Statistics* 29 (3): 466–78. <https://doi.org/10.1080/10618600.2019.1695624>.
- Wickham, Hadley. 2011. "The Split-Apply-Combine Strategy for Data Analysis." *Journal of Statistical Software, Articles* 40 (1): 1–29. <https://doi.org/10.18637/jss.v040.i01>.
- . 2014. "Tidy Data." *Journal of Statistical Software* 59 (10): 1–23.
- . 2016. *Ggplot2: Elegant Graphics for Data Analysis*. Springer-Verlag New York. <https://ggplot2.tidyverse.org>.
- . 2019. *Stringr: Simple, Consistent Wrappers for Common String Operations*. <https://CRAN.R-project.org/package=stringr>.
- . 2020a. *Modelr: Modelling Functions That Work with the Pipe*. <https://CRAN.R-project.org/package=modelr>.
- . 2020b. *Tidyr: Tidy Messy Data*. <https://CRAN.R-project.org/package=tidyr>.
- Wickham, Hadley, Mara Averick, Jennifer Bryan, Winston Chang, Lucy D'Agostino McGowan, Romain François, Garrett Grolemond, et al. 2019. "Welcome to the tidyverse." *Journal of Open Source Software* 4 (43): 1686. <https://doi.org/10.21105/joss.01686>.
- Wickham, Hadley, Romain François, Lionel Henry, and Kirill Müller. 2020. *Dplyr: A Grammar of Data Manipulation*. <https://CRAN.R-project.org/package=dplyr>.
- Wickham, Hadley, and Jim Hester. 2020. *Readr: Read Rectangular Text Data*. <https://CRAN.R-project.org/package=readr>.
- Wood, S. N. 2003. "Thin-Plate Regression Splines." *Journal of the Royal Statistical Society (B)* 65 (1): 95–114.
- Xie, Yihui. 2014. "Knitr: A Comprehensive Tool for Reproducible Research in R." In *Implementing Reproducible Computational Research*, edited by Victoria Stodden, Friedrich Leisch, and Roger D. Peng. Chapman; Hall/CRC. <http://www.crcpress.com/product/isbn/9781466561595>.
- Xie, Yihui, Christophe Dervieux, and Emily Riederer. 2020. *R Markdown Cookbook*. Boca Raton, Florida: Chapman; Hall/CRC. <https://bookdown.org/yihui/rmarkdown-cookbook>.

Zhu, Hao. 2019. *kableExtra: Construct Complex Table with 'Kable' and Pipe Syntax*. <https://CRAN.R-project.org/package=kableExtra>.